

Sami Habbal

Charlotte, NC · samihabbal5@icloud.com · gg5533.github.io · github.com/GG5533

I work on the failures that only appear under real traffic: retries that cause double effects, provider calls with no deadline, agent loops with no ceiling, and state that disagrees with what the payment provider thinks happened.

OPEN-SOURCE RELIABILITY WORK

changedetection.io — LLM client retry semantics

33k★ · Sept 2026

- Reviewed the retry handling in the LLM client. The maintainer replied "*These points are spot on*" and implemented the findings in [PR #4384](#) — `Retry-After` per RFC 9110, backoff jitter, and a delay ceiling. PR open, not yet merged.
- Wrote regression tests for the retry-budget behaviour that fix introduced; **the maintainer merged them into that branch** (commit `1f0ed3f`, now one of the four commits in PR #4384). A stripped sampling-parameter retry refunds its attempt, so `DEFAULT_RETRIES + 1` provider calls are reachable — verified by mutation, since deleting `attempt -= 1` fails that test and no other in the file.
- Filed [#4385](#): the AI intent filter fails *closed* on a malformed model response — suppressing the change and caching that verdict — while every other failure path in the same function fails open.
- Reported a private security advisory (GHSA-pv58-w252-g26w): an unescaped LLM-derived variable reaching HTML notification bodies, the same class as their existing GHSA-q8xq-qg4x-wphg.

Circle · arc-node — SIGTERM shutdown race

Rust · Sept 2026

- Reviewed a fix for a shutdown race in their consensus node. The fix guards the path where the app task resolves normally, but the `?` on the line above the new check discharges a `JoinError` first — so a panicking task still drops the runtime and kills cleanup mid-flight.
- Independently confirmed by another contributor against a standalone tokio model: "@GG5533's `?` *finding is real*."
- Corrected my own scope publicly when a reviewer showed the durability consequence was narrower than I first claimed. [circlefin/arc-node#361](#)

SELECTED PROJECTS

stripe-fulfilment-kit — fenced leases for order fulfilment

TypeScript · [repo](#)

- A lease implementation that passed 78 tests and was still wrong: a provider call outliving its lease could write failure over a completed order, hiding a paid-for product and reopening it for a duplicate charge. Every test settled an order from its *current* owner; none covered a write from a superseded one.
- Fixed with fencing tokens enforced atomically across processes via `BEGIN IMMEDIATE`, plus regression tests for the case the original suite could not see. [Full write-up](#).

ai-review-gate — adversarial review before output ships

[repo](#)

- Two models review each substantive answer in parallel and block it if either finds a real defect. It has caught claims that outran their evidence — including one asserting a guard prevented an unbounded loop, which I disproved by deleting the guard and watching every test still pass.

TECHNICAL

TypeScript, Node, Python, Rust (reading/review), Next.js, Stripe, SQLite/Postgres, LLM APIs (OpenAI/Anthropic), agent orchestration, distributed-systems debugging, formal reasoning about concurrency (leases, fencing, idempotency).

Every claim above links to a public artifact. Merge status is stated where a contribution is not yet merged.